

Presentation 5 - Celena Kim

Presentation Links:

- No Power BI Dashboard link- was simply tasked with calculating the time it takes for callers to move through the menu queues before and after the Spring Quarter 2025 implementation.
- Code
 - Can also refer to the first part of [Jake's Dec 9th code](#)
- [Final Presentation Slides](#)

What did you do:

- For this final presentation analysis, I focused on Objective 2A of validating whether the Spring Quarter recommendation implemented on October 16, 2025 of removing the "Pre-legal seniors menu" successfully reduced the time it takes for a caller to reach the next menu.
- In order to verify this improvement, I performed the following tasks:
 - Augmented our current CAR dataset with the new October/November call data provided by Cynthia
 - Separated calls into two groups:
 - Pre-implementation: all calls before October 16, 2025
 - Post-implementation: all calls on or after October 16, 2025
 - Then, for each call session in the pre and post implementation sets, I calculated the time it took for a caller to go from SeniorsMenu to SuburbsOrCityMenu.
 - Our expectation was that traversal times would be faster in the post-implementation data, since callers would no longer be routed through an additional, unnecessary menu.

How does it help the project?:

- The results clearly show that the Spring Quarter recommendation had a meaningful impact.
 - **Before implementation:** it took callers an average of **23.7 seconds** to navigate from SeniorsMenu to SuburbsOrCityMenu.
 - **After implementation:** this decreased to **7.1 seconds**.
- This nearly 60% reduction in navigation time demonstrates that this Spring Quarter implementation significantly reduces the time it takes for a caller to flow through the menus, resulting in a substantially more efficient and improved caller experience.

Issues faced (if any):

- None

Attempts to resolve issues

- N/A

Presentation 4 - Celena Kim

Presentation Links:

- [Power BI Dashboard Link](#) (was not able to convert to .pbix when attempted?)
 - Updated 11/19 5:35 PM– realized I was not using my updated dataset
- [Termination Reason Analysis Code](#)
- [Presentation 4 Slides](#) (Just my section)
- [Presentation 4 Recording](#) (Google Drive link)

What did you do:

- Initially, our group's objective was to run regression models to create 95% prediction intervals, but given the poor model results, we decided to shift focus. While examining different call sessions in the CARs dataset, I noticed irregularities within several of the termination patterns. Ideally, a call should end at the first row that contains a termination reason such as "Customer Left". However, I identified several anomalies such as sessions continuing long after a termination event or containing multiple termination reasons (e.g. "Customer Left" followed by "Agent Left"). These irregularities indicate potential inconsistencies in how the call platform is logging events, so I wanted to explore this further. To investigate, I defined the following logic for classifying a "Normal" vs. "Irregular" sessions:

Normal session:

- Only one termination reason that is not repeated
- $\text{first_termination_time} - \text{last_activity_time} = 0$
- Multiple termination reasons only if the final one is MAX_CALLBACK_RETRY_LIMIT_REACHED, and prior ones are expected retries like CUSTOMER_UNAVAILABLE.

Irregular session:

- Multiple termination reasons/Termination events not in the expected order (e.g., Customer Left then Agent Left)
- Significant amount of elapsed time after first termination event
- I then visualized the distribution of normal vs. irregular calls with a bar chart within Power BI and found that around 5% of the calls had irregular terminations. To enable further exploration, I also implemented a drill-through feature so that users can search for irregular calls and view the full session by contact session ID, helping identify what occurred and what was incorrectly logged.
- ***NOTE:**
- 2 datasets are used within this analysis:
 - clean_session_df: all the rows in the CARs dataset

- TerminationReasonAnalysis_18Nov: a summary of termination details for each contact_session_id, including termination reason count and whether it's classified as "Normal" or "Irregular"
- Within Power BI, there is a many to 1 relationship from clean_session_df to TerminationReasonAnalysis_18Nov, joined by contact_session_id

How does it help the project?:

- This analysis reveals a potential data logging inconsistency in the CARs dataset, specifically around how call session terminations are recorded. By identifying and flagging these anomalies, we can establish a more accurate logic for logging call sessions by ensuring that sessions end at the first instance of a termination, rather than lingering on the line and potentially taking time away from getting back to other callers waiting for a callback or picking up incoming calls. This not only improves the quality and reliability of the dataset for future analyses, but also highlights operational inefficiencies in call recording. Accurately capturing session endings can help reduce unnecessary agent idle time and ensure that callbacks and incoming calls are handled promptly.

Issues faced (if any):

- None

Attempts to resolve issues

- N/A

Next Steps

- I plan to explore the impact of these anomalies on agent performance metrics and customer experience in order to assess whether session prolongation following the multiple recorded terminations are affecting operational efficiency. For example, I could assess patterns across specific agents and call types to identify whether certain conditions or individuals are driving these anomalies, which could inform targeted process improvements or adjustments to the call platform's logging system.

Presentation 3 - Celena Kim

Presentation Link:

- [Power BI Dashboard](#)
- [Callback & Post-Queue Pickup Analysis Code](#)
- [Presentation 3 Slides](#)

What did you do:

- For this presentation, I joined Jake and Logan in analyzing the CAR dataset. Building off of Jake and Logan's analyses, which mainly focused on queue wait times and callbacks vs. no callbacks, I wanted to focus on analyzing what happens after a customer connects with an agent. Specifically, I examined the duration of calls for individuals who used the callback feature versus those who stayed in the queue to wait for the agent. I defined a session as utilizing the callback feature if it contained at least one activity where `flow_name == "CourtesyCallback"`. For these sessions, I determined whether the callback was successful if they eventually reached an agent (indicated by the last instance of "LegalServerScreenPop" in the `activity_name` column) and if the rows following this point did not immediately contain a termination reason such as "Customer Left", "Agent Left", "No Answer From Agent", etc. I then calculated the call duration from the start time of the last LegalServerScreenPop activity to the last row in that call session. For the post-queue pickup analysis, I identified sessions that did not request callbacks but included queue-related activities such as PreQueue messages or Queue menus. In a similar fashion, call duration was calculated as, following a queue, the time elapsed between when the session reached an agent (the last LegalServerScreenPop timestamp) and the end of the session. After calculating the call durations for these two call types, I then visualized two graphs within Power BI: the overall number of calls per hour by call type, and the average duration of calls per hour by call type. I also included cards displaying the average duration of each call type, as well as a date slicer to filter by different periods.

How does it help the project?:

- This analysis provides valuable insights into comparing how calls are handled as a callback vs. post-queue pickups. The results show that callbacks not only are significantly more prevalent but also slightly longer in duration, suggesting that individuals strongly prefer to utilize the callback feature rather than wait on hold in a queue. Moreover, agents may either be spending more focused time with individuals during callbacks, leading to the slightly longer average durations, or should streamline their processes in order to decrease these call durations. These insights can inform call

optimization by prioritizing resources and training towards handling these callbacks more efficiently since they represent the majority of customer interactions. Furthermore, visualizing these trends by the hour reveals that peak call times for both callback and post-queue pickups are 8 AM, 12 PM, and 4 PM, suggesting that adequate staffing and workflow adjustments may be necessary to manage these high-traffic periods effectively. Overall, this insight into agent call handling highlights clear patterns and opportunities to improve caller experience by focusing on timing, prevalence, and management of callbacks vs. post-queue calls.

Issues faced (if any):

- One issue I faced was not being too sure if my logic for defining successful callbacks/post-queue pickups and calculating call durations was correct. There were some distinct cases such as a call session containing multiple instances of “LegalServerScreenPop” that I was unsure how to interpret.
- For instance, contact session id “01a50ccf-91f5-403c-92ae-1b77bba1dff4” had three instances of “LegalServerScreenPop” across two days, with two different agents, Ana Franco and Marisol Guadarrama. I was unsure whether all three of these callback instances should have been considered as 3 distinct callback sessions.

Attempts to resolve issues

- To resolve this uncertainty, I set a consistent approach to just consider the last instance of “LegalServerScreenPop”, coupled with the absence of a termination reason, to represent a successful callback or post-queue pick up. I also calculated call duration as the point of the last instance of “LegalServerScreenPop” to the last row of the call session.

Next Steps

- I plan to revisit my logic for defining callback/post-queue pickup success by attempting to handle multiple callback instances such as what we saw with contact session id “01a50ccf-91f5-403c-92ae-1b77bba1dff4”. Once this logic is clarified, I would like to focus in on this analysis in more detail, such as adding in a filter for different queue types.

Presentation 2 - Celena Kim

Presentation Link:

- [Power BI Dashboard](#)
- [Call Duration Data Prep Code](#)
- [Presentation 2 Slides](#)

What did you do:

- For this analysis, I used the All Calls dataset to visualize call counts and durations for distinct called numbers. I first defined Legal Aid's open and closed hours and created a reference table to map phone number descriptions from the "Popular Numbers (AllCallsData)" spreadsheet. I then removed NA called numbers and any duplicate rows that corresponded to the same call log, which reduced the dataset from 1,001,537 rows to 732,420 rows. Next, to enable an interactive filtering feature in the dashboard, I created several additional columns: call path (with values inbound, outbound, and internal, derived using the call direction and PSTN name), date, weekday vs. weekend, and open vs. closed hours. This prepared dataset allowed me to graph the distribution of call count and duration and categorize the calls by multiple dimensions of filters, enabling deeper insights into call patterns at both working and non-working hours.

How does it help the project?:

- By exploring call counts and durations across numerous categories of call path, day of the weekend, and open vs. closed hours, this analysis helps the project by identifying any trends or bottlenecks in service demand, highlighting peak periods or unnecessarily long call durations at various points, and enabling the Legal Aid to make any decisions to improve call response times and efficiencies. Specifically, the dashboard revealed that during closed hours on weekdays, the Farmworker main number (Migrant Legal Assistance Program) line displayed significantly high average durations of 1,182 minutes, suggesting a need to evaluate the process implemented for handling calls outside working hours and potentially implement some sort of automated system that immediately informs callers that the office is closed and points them to online resources, rather than letting those calls linger on the line. Additionally, on weekends when the office is completely closed, both the Farmworker main number as well as the "Transfer to the English main menu" lines exhibited unusually high average durations of around 1,282 minutes and 1,859 minutes, respectively, indicating that there may be a potential bottleneck within these lines on weekends with callers either waiting excessively or somehow reaching the unstaffed line. For this potential issue, a revised automated

messaging process could also reduce this line's high duration and provide alternative resources to these weekend callers.

Issues faced (if any):

- One issue I faced is not being too sure that I prepped my data correctly. Based on the feedback from project 1, rather than grouping by call direction (Terminating/originating), I created a column called "Call Path" with values inbound, outbound, or internal. This was derived by using the call direction and PSTN name, and utilized within the dashboard as a filter. Next, in order to not overcount duplicate logs of the same call, I removed any duplicate rows corresponding to the same combination of `Correlation ID` and `Called number`, and same combination of `Correlation ID`, `Start time`, and Duration. However, this reduced the dataset from 1,001,537 rows to 732,420 rows, meaning some analyses were limited to smaller subsets. Given that my last recommendation from the previous presentation was to avoid analyses that only focus on a small subset of calls, I attempted to produce results that focused on higher proportions of the dataset, however, I still found it difficult to extract meaningful insights from larger subsets, since clear bottlenecks were found in areas that represented 11% of the dataset. Additionally, many of the phone numbers did not have descriptions, so it was difficult to interpret some of my results.

Attempts to resolve issues/Next Steps

- In order to address producing insights for larger portions of the dataset, plan to analyze more significant categories for analysis. Specifically, since I was able to identify that the main line called number category clearly holds the largest proportion of calls from this analysis, I could focus on this subset and explore the different pathways that extend from the main line to identify any potential bottlenecks. For instance, perhaps a pathway for a caller to get to a certain line after being funneled through the main line is exorbitantly long-winded and tedious. By zooming in on a known significant portion of the dataset, I am more likely to uncover actionable patterns that can reflect the broader call landscape, rather than pinpointing minor improvements.

Presentation 1 - Celena Kim

Presentation Link:

- [Power BI Dashboard](#)
- No extra code files– DAX measures and new columns were created within Power BI:
 - Avg Duration = AVERAGE(all_calls[Duration])
 - Call_Column = 'all_calls'[Direction] & " | " & 'all_calls'[Releasing party] & " | " & 'all_calls'[Call outcome]

What did you do:

- To identify potential inefficiencies within Legal Aid's call process, I analyzed patterns in average call duration across various call combinations. Using the All Calls dataset, I created a DAX measure in Power BI to calculate the average durations for each unique combination of Direction (Terminating in Legal Aid or Originating), Releasing Party (Local, Remote, or Unknown), and Outcome (Refusal, Failure, or Success). I then visualized these results in a graph to display this distribution of durations. Finally, to make key patterns more visible, I applied conditional color formatting to highlight combinations with notably higher durations in red. Three distinct outliers emerged from this analysis, all with a common trait: having an "Unknown" releasing party.

How does it help the project?:

- This finding highlights a clear opportunity to optimize how the system manages calls with an "Unknown" releasing party. Since these calls tend to have significantly longer average durations, Legal Aid could improve their efficiency in various ways such as by automatically terminating unknown caller IDs, implementing a timeout mechanism for these unknown calls instead of letting them linger on the line, or further investigating why certain callers appear as unknown in order to allow these callers to still receive the help they deserve. These solutions would not only prevent wasted call times, but also allow for legitimate callers with these unknown caller IDs to still receive assistance.

Issues faced (if any):

- After displaying the counts of these 3 "Unknown" call combinations with the bar chart on the right side of the dashboard, I found that these instances are very few in number within the vast 1 million rows of calls in the dataset, with counts only ranging from 2-158. Therefore, while these outliers do highlight an interesting inefficiency, the overall impact of this issue is minimal in scope. This limitation suggests that while the dashboard does

flag a bottleneck within the call process, further analysis is necessary to identify an inefficiency that has a larger impact on overall performance.

Attempts to resolve issues/Next Steps

- While I am still interested in this pattern of lingering call durations, I plan to identify more widespread patterns that have a greater impact within the scope of this large dataset. For instance, I could attempt to analyze different call combinations with higher frequencies to uncover duration inefficiencies, or explore other areas for optimization besides such as agent response times. Overall, I aim to gain a more comprehensive understanding of where the system could be improved by highlighting prolonged interactions or delays that may be reducing call efficiency, while keeping in mind to identify business insights that have a meaningful effect at scale.